

Fuzzing the CNCF landscape

Adam Korczynski, Ada Logics <adam@adalogics.com>
David Korczynski, Ada Logics <david@adalogics.com>

ADALOGICS



CLOUD NATIVE
COMPUTING FOUNDATION

Agenda

- Overview
- Fuzzing quick intro
- Open source fuzzing with OSS-Fuzz
- The CNCF landscape being fuzzed
- How to fuzz a CNCF project
- Fuzzing results
- Future work

Fuzzing: quick intro

Testing versus fuzzing

Test:

```
MyApi(Input1);  
MyApi(Input2);  
MyApi(input3);
```

Fuzzing:

```
while true {  
    MyApi(Fuzzer.GenerateInput());  
}
```

Testing versus fuzzing

Test:

```
MyApi(Input1);  
MyApi(Input2);  
MyApi(input3);
```

Fuzzing:

```
while true {  
    MyApi(Fuzzer.GenerateInput());  
}
```

Fuzzing - actual implementation:

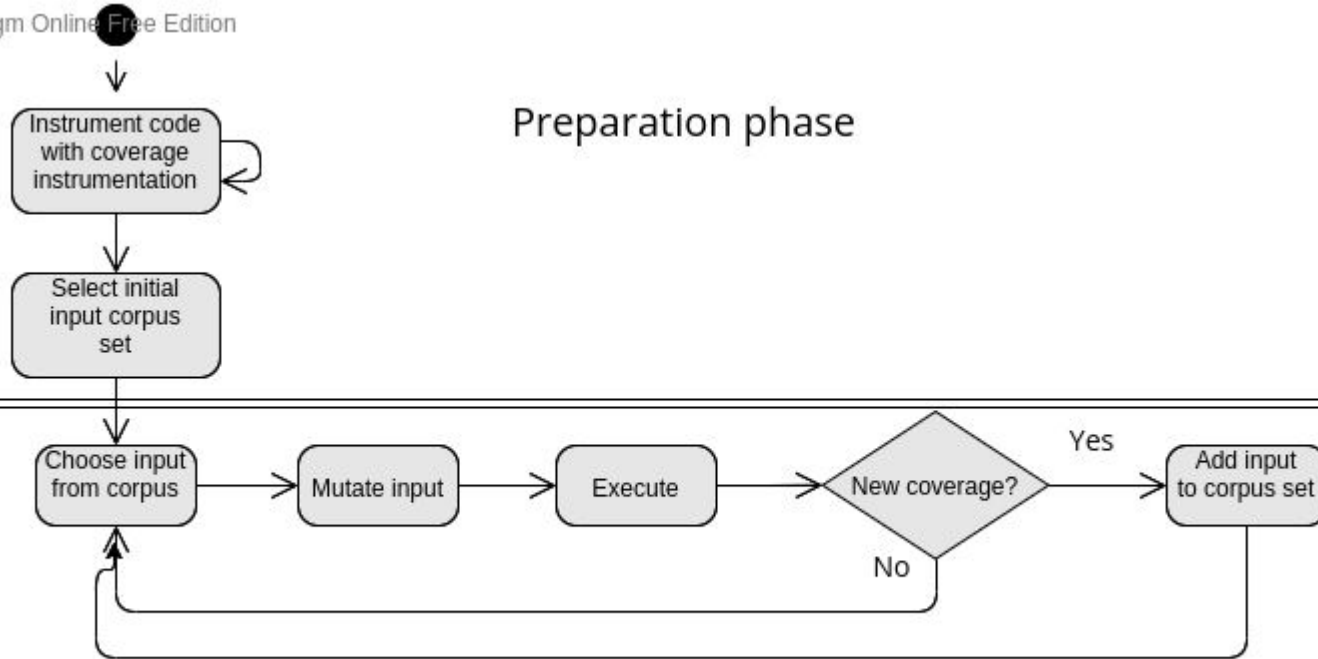
```
func Fuzz(data []byte) int {  
    MyApi(string(data))  
    return 1  
}
```

Fuzzing algorithmic underpinnings

- Common myth: “Fuzzing is random testing, it will never analyse complex code”
 - **False!**
- The origins of fuzzing is in random testing
- Modern day fuzzers are complex genetic mutational algorithms
- Mutations involve a random element
- “Modern day fuzzers” in this case refer to *coverage guided fuzzers*
- Hundreds of academic papers on how to improve fuzzing engines

Coverage guided fuzzing

Visual Paradigm Online Free Edition



Execution phase


Visual Paradigm Online Free Edition

Intuition for coverage-guided fuzzing

Current seed: ____

Probability of progression: $1/256$

Probability of no progression: $255/256$




Intuition for coverage-guided fuzzing

Current seed: A__

Probability of progression: $1/256$

Probability of no progression: $255/256$




Intuition for coverage-guided fuzzing

Current seed: AB__

Probability of progression: $1/256$

Probability of no progression: $255/256$




Intuition for coverage-guided fuzzing

Current seed: ABC_

Probability of progression: $1/256$


Probability of no progression: $255/256$



Intuition for coverage-guided fuzzing

Current seed: ABCD

Bug found after $256 * 4$ tries



Which bugs can fuzzers uncover?

- Fuzzers are conceptually only test-case generators. They generate inputs that trigger code paths, i.e. they are not bug identifiers as such but rather execution path identifiers.
- **Memory unsafe languages**
 - Memory corruption.
 - Everything sanitizers tell you
 - ...
- **Memory safe languages**
 - Uncaught exceptions
 - Out-of-bounds
 - Nil-pointer dereferences
 - Time outs
 - Out of memory issues
- Behavioural testing

Open source fuzzing with OSS-Fuzz

OSS-Fuzz

- The management of fuzzers is a complex task
- OSS-Fuzz is a free service managed through Github that:
 - Runs fuzzers in open source projects
 - Filters and analyses data from the fuzzers
 - Reports when issues are found to maintainers
 - Tracks when bugs are fixed
 - Suggests improvements and fuzz introspection capabilities
 - <https://github.com/google/oss-fuzz>
- Integration into OSS-Fuzz is easy
 - “Integrating fuzzing into your open source project with OSS-Fuzz available [here](#)

ClusterFuzzLite


- ClusterFuzzLite is an extension to Clusterfuzz which OSS-Fuzz uses
- Runs as part of the CI
- Catches bug early in the process
- Projects can (also) integrate this to catch bugs in the CI, and also integrate into ClusterFuzzLite without full integration into OSS-Fuzz
- <https://google.github.io/clusterfuzzlite/>

OSS-Fuzz project integration

1. Develop one or more fuzzers for a project
2. Develop a Dockerfile that builds the environment in which the project and fuzzers can be built
3. Develop a build.sh script that builds all the fuzzers
4. Create a project.yaml file with maintainer emails
5. Make a pull request on <https://github.com/google/oss-fuzz>
6. Once the PR is merged, OSS-Fuzz will run the fuzzers indefinitely

The CNCF landscape being fuzzed

CNCF projects being fuzzed



How to fuzz a CNCF project

High level process when fuzzing a CNCF project

High level process when fuzzing a CNCF project

1: Initial integration

High level process when fuzzing a CNCF project

1: Initial integration



High level process when fuzzing a CNCF project

1: Initial integration



High level process when fuzzing a CNCF project

1: Initial integration



2: Write a lot of fuzzers

High level process when fuzzing a CNCF project

1: Initial integration



2: Write a lot of fuzzers

3: Additional goals depending on project

Integrate as a third-party

- As a third party we bring a lot of knowledge about fuzzing, but have little knowledge of the target project
- Maintainers have a lot of knowledge about target project, but little knowledge of fuzzing
- Maintainers have little time available

Encourage and help maintainers as third parties

- Integrate the fuzzers into the upstream repository
- Do root-cause analysis of bugs reported
- Write more fuzzers
- Inspect the state of the fuzzing to identify limitations
- Take complete control of the fuzzing process

Go-fuzz-headers

<https://github.com/AdaLogics/go-fuzz-headers>

```
import (  
    fuzz "github.com/AdaLogics/go-fuzz-headers"  
)  
func Fuzz(data []byte) int {  
    f := fuzz.NewConsumer(data)  
  
    myStruct := &CustomStruct{}  
    err := f.GenerateStruct(myStruct)  
  
    myMap := make(map[string]string)  
    err := f.FuzzMap(&myMap)  
  
    var mySlice []string  
    err := f.CreateSlice(&mySlice)  
    return 1  
}
```

Go-fuzz-headers

<https://github.com/AdaLogics/go-fuzz-headers>

```
import (  
    fuzz "github.com/AdaLogics/go-fuzz-headers"  
)  
func Fuzz(data []byte) int {  
    f := fuzz.NewConsumer(data)  
  
    myStruct := &CustomStruct{}  
    err := f.GenerateStruct(myStruct)  
  
    myMap := make(map[string]string)  
    err := f.FuzzMap(&myMap)  
  
    var mySlice []string  
    err := f.CreateSlice(&mySlice)  
    return 1  
}
```

Go-fuzz-headers

<https://github.com/AdaLogics/go-fuzz-headers>

```
import (  
    fuzz "github.com/AdaLogics/go-fuzz-headers"  
)  
func Fuzz(data []byte) int {  
    f := fuzz.NewConsumer(data)  
  
    myStruct := &CustomStruct{}  
    err := f.GenerateStruct(myStruct)  
  
    myMap := make(map[string]string)  
    err := f.FuzzMap(&myMap)  
  
    var mySlice []string  
    err := f.CreateSlice(&mySlice)  
    return 1  
}
```

Go-fuzz-headers

<https://github.com/AdaLogics/go-fuzz-headers>

```
import (  
    fuzz "github.com/AdaLogics/go-fuzz-headers"  
)  
func Fuzz(data []byte) int {  
    f := fuzz.NewConsumer(data)  
  
    myStruct := &CustomStruct{}  
    err := f.GenerateStruct(myStruct)  
  
    myMap := make(map[string]string)  
    err := f.FuzzMap(&myMap)  
  
    var mySlice []string  
    err := f.CreateSlice(&mySlice)  
    return 1  
}
```


Go-fuzz-headers

<https://github.com/AdaLogics/go-fuzz-headers>

```
import (  
    fuzz "github.com/AdaLogics/go-fuzz-headers"  
)  
func Fuzz(data []byte) int {  
    f := fuzz.NewConsumer(data)  
  
    myStruct := &CustomStruct{}  
    err := f.GenerateStruct(myStruct)  
  
    myMap := make(map[string]string)  
    err := f.FuzzMap(&myMap)  
  
    var mySlice []string  
    err := f.CreateSlice(&mySlice)  
    return 1  
}
```

The spectrum of fuzzing CNCF projects

- There are CNCF projects in all sorts of languages
- Fuzzing by way of OSS-Fuzz can be done in the languages
 - Golang
 - C/C++
 - Rust
 - Python
 - Java

Bugs to find in CNCF projects

- The bugs you find largely depends on the language of the project
- Threat models are not always present
- CNCF projects are often written in memory safe languages
- Dependencies

Bugs to find in CNCF projects

- The bugs you find largely depends on the language of the project
- Threat models are not always present
- CNCF projects are often written in memory safe languages
- Dependencies



Bugs to find in CNCF projects

- The bugs you find largely depends on the language of the project
- Threat models are not always present
- CNCF projects are often written in memory safe languages
- Dependencies



Bugs to find in CNCF projects

- The bugs you find largely depends on the language of the project
- Threat models are not always present
- CNCF projects are often written in memory safe languages
- Dependencies



flux



git2go



libgit2



60k loc

Fuzzing results

Querying monorail for issues

- URL: <https://bugs.chromium.org/p/oss-fuzz/issues/>
 - Select “All issues” in drop-down list.
- Query: “proj=PROJ_NAME Type=Bug-Security,Bug label:Reproducible”
- This count gives over-approximations on what is considered real bugs.

Fuzzing results

Project	Approx bugs found	Language
Argo	25	Go
Cluster API	4	Go
Containerd	4	Go
CoreDNS	9	Go
CRI-O	4	Go
Distribution	4	Go
Envoy	869	C++
Etcd	15	Go
Fluent-Bit	222	C
Helm	2	Go
Kubernetes	54	Go
Linkerd2 + (-proxy)	7 + (14)	Go + Rust
Runc	0	Go
Vitess	45	Go

Fuzzing results

Project	Approx bugs found	Language
Argo	25	Go
Cluster API	4	Go
Containerd	4	Go
CoreDNS	9	Go
CRI-O	4	Go
Distribution	4	Go
Envoy	869	C++
Etcd	15	Go
Fluent-Bit	222	C
Helm	2	Go
Kubernetes	54	Go
Linkerd2 + (-proxy)	7 + (14)	Go + Rust
Runc	0	Go
Vitess	45	Go

Fuzzing results

Project	Approx bugs found	Language
Argo	25	Go
Cluster API	4	Go
Containerd	4	Go
CoreDNS	9	Go
CRI-O	4	Go
Distribution	4	Go
Envoy	869	C++
Etcd	15	Go
Fluent-Bit	222	C
Helm	2	Go
Kubernetes	54	Go
Linkerd2 + (-proxy)	7 + (14)	Go + Rust
Runc	0	Go
Vitess	45	Go


Plotting issues

- Plotting issues from monorail show macro effects of fuzzing
 - Use the “chart” feature on monorail in the previous queries
- Closed issues: bugs reported by OSS-Fuzz that are now closed
- Open issues: bugs reported by OSS-Fuzz that are not resolved.

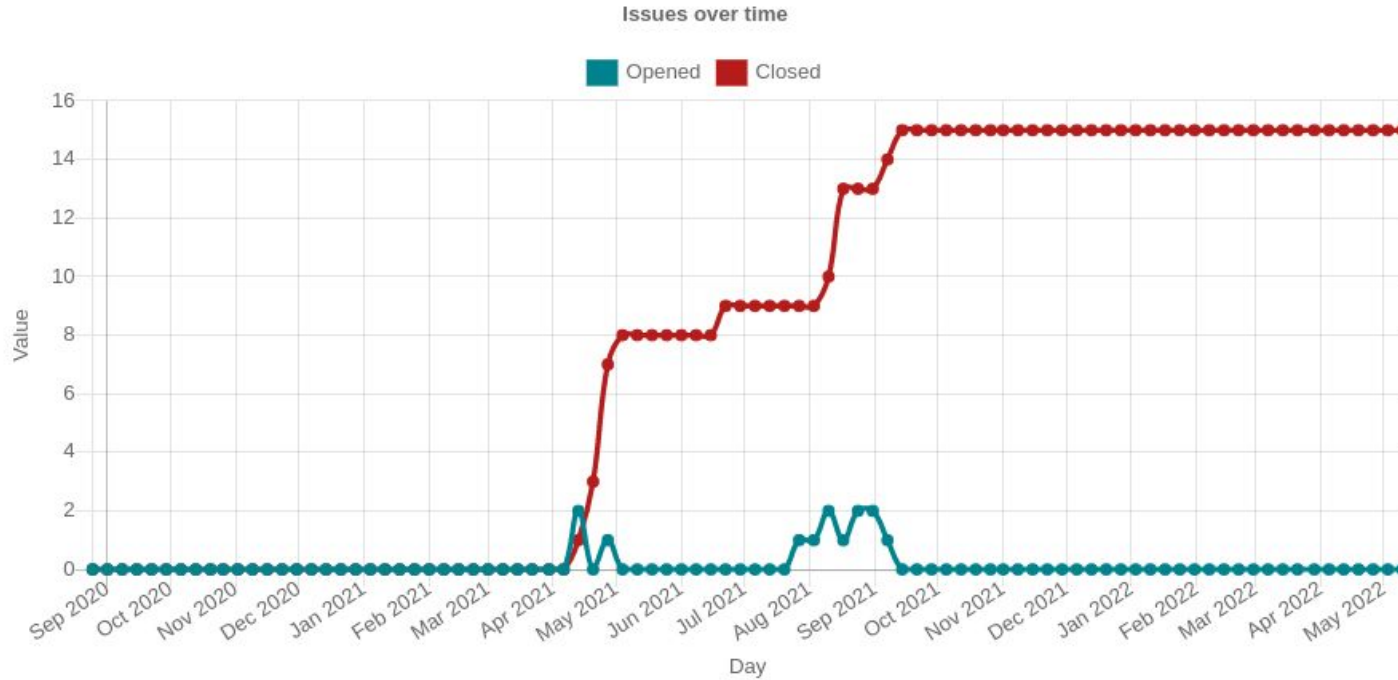
Argo




Kubernetes



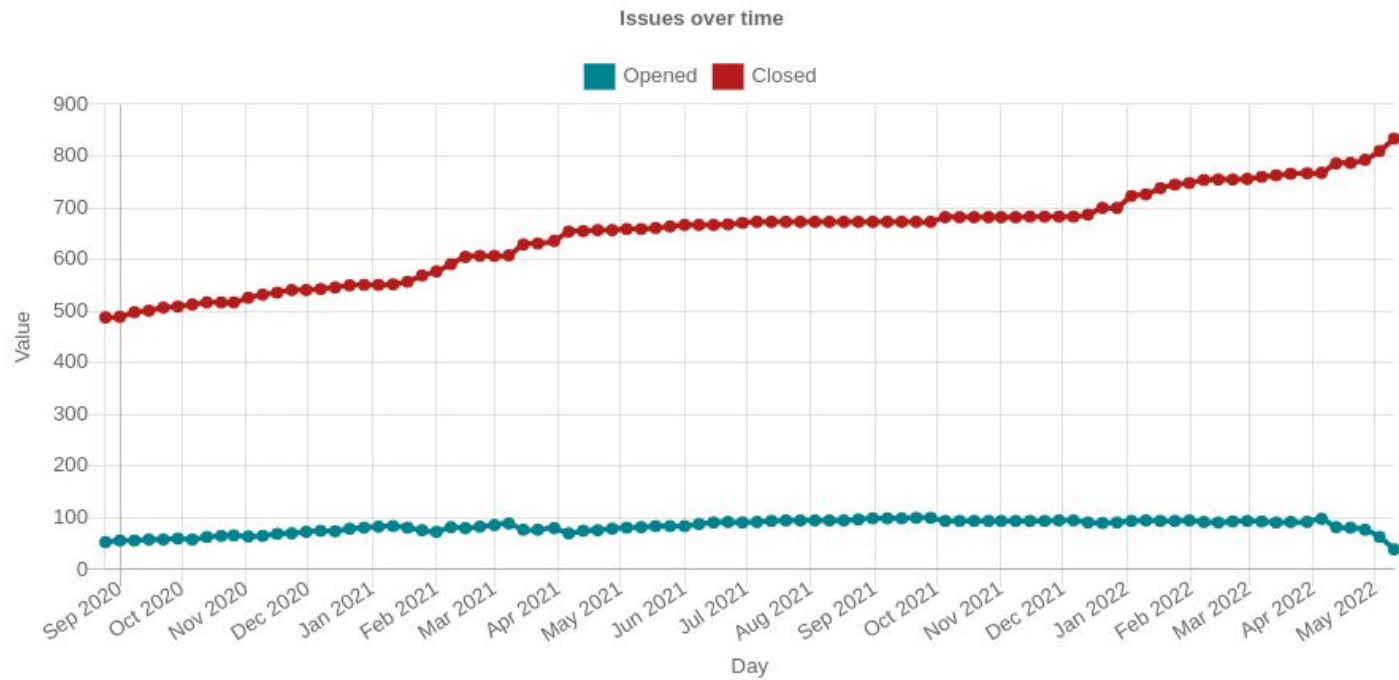
Linkerd2-proxy



Fluent Bit

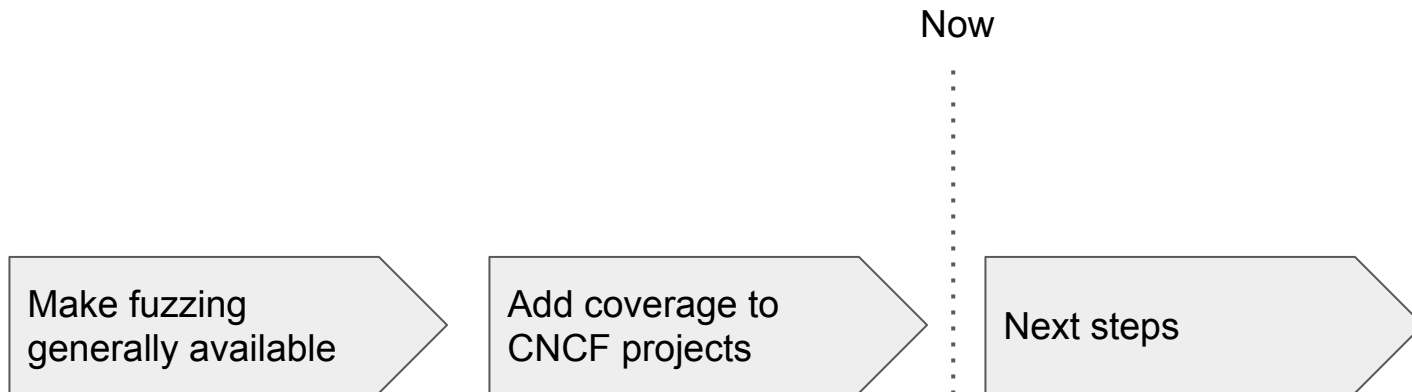


Envoy

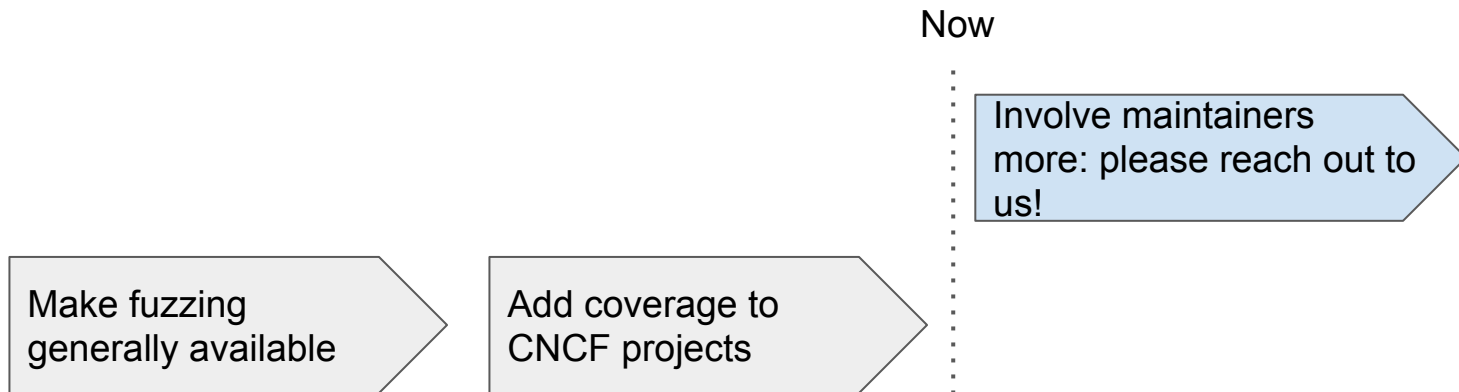


Future work

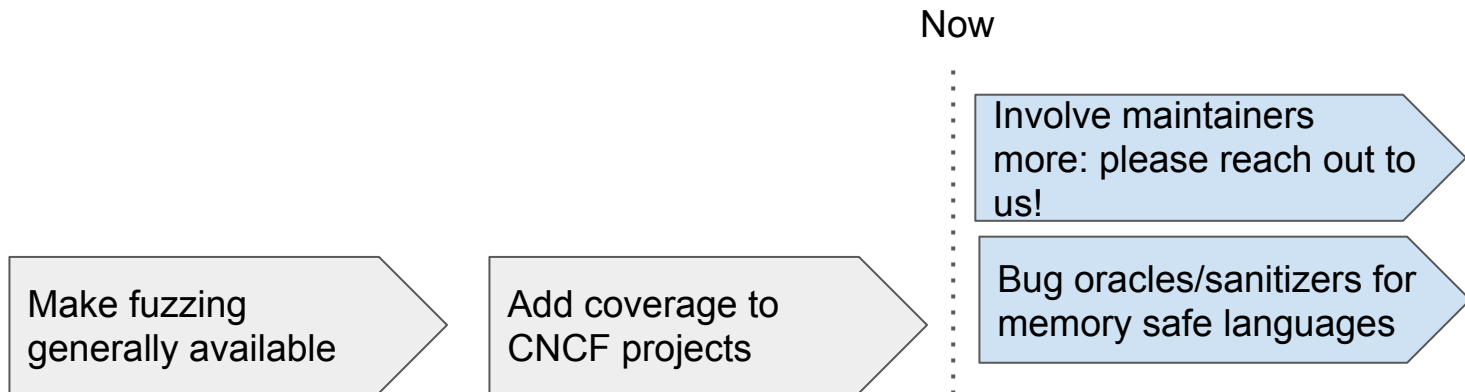
Looking forward



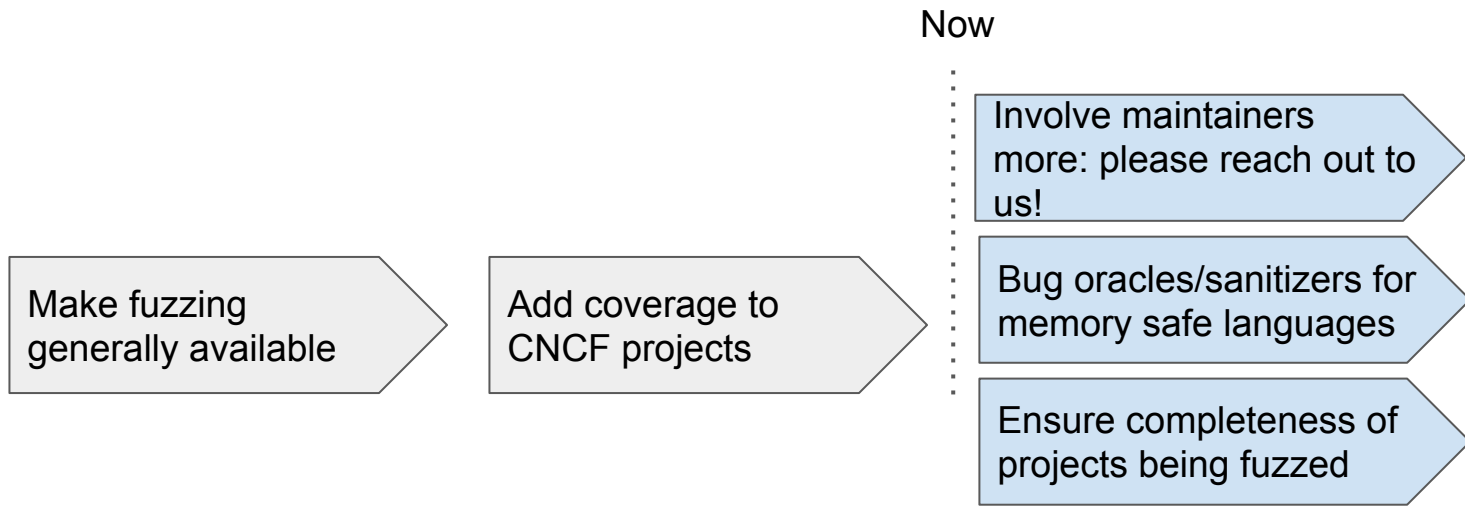
Looking forward



Looking forward



Looking forward



Increase maintainer involvement

- Showcase importance of fuzzing through results
- Assist maintainers in fuzzing development/integration
- Provide documentation and training material
- Please reach out via <https://github.com/cncf/cncf-fuzzing> if you'd like your project fuzzed

Better bug detection

- Applied in a general manner.

Better bug detection

- Applied in a general manner.
- Go fuzzing:
 - Out of range
 - Out of memory
 - Time outs

Better bug detection

- Applied in a general manner.
- Go fuzzing:
 - Out of range
 - Out of memory
 - Time outs
 - **Race conditions**

Better bug detection

- Applied in a general manner.
- Go fuzzing:
 - Out of range
 - Out of memory
 - Time outs
 - **Race conditions**
 - **Logging**

Better bug detection

- Applied in a general manner.
- Go fuzzing:
 - Out of range
 - Out of memory
 - Time outs
 - **Race conditions**
 - **Logging**
 - **File handling**

Better bug detection

- Applied in a general manner.
- Go fuzzing:
 - Out of range
 - Out of memory
 - Time outs
 - **Race conditions**
 - **Logging**
 - **File handling**
 - **Command injections**

Better bug detection

- Applied in a general manner.
- Go fuzzing:
 - Out of range
 - Out of memory
 - Time outs
 - **Race conditions**
 - **Logging**
 - **File handling**
 - **Command injections**
- Massive impact



Conclusions

Conclusions

- 15+ CNCF projects has integrated fuzzing over the last two years
- Integrating with OSS-Fuzz provides a continuous security monitoring on each CNCF project
- Many bugs have been discovered
- Security and reliability bugs
- The fuzzing work has several future works:
 - Extend fuzzers
 - New bug oracles for memory safe languages
 - Involve maintainers more